Code and Data for the Social Sciences:
A Practitioner's Guide

Matthew Gentzkow   Jesse M. Shapiro[1]
Chicago Booth and NBER

March 10, 2014

Contents

Chapter 1

Introduction

What does it mean to do empirical social science? Asking good questions. Digging up novel data. Designing statistical analysis. Writing up results.

For many of us, most of the time, what it means is writing and debugging code. We write code to clean data, to transform data, to scrape data, and to merge data. We write code to execute statistical analyses, to simulate models, to format results, to produce plots. We stare at, puzzle over, fight with, and curse at code that isn't working the way we expect it to. We dig through old code trying to figure out what we were thinking when we wrote it, or why we're getting a different result from the one we got the week before.

Even researchers lucky enough to have graduate students or research assistants who write code for them still spend a significant amount of time reviewing code, instructing on coding style, or fixing broken code.

Though we all write code for a living, few of the economists, political scientists, psychologists, sociologists, or other empirical researchers we know have any formal training in computer science. Most of them picked up the basics of programming without much effort, and have never given it much thought since. Saying they should spend more time thinking about the way they write code would be like telling a novelist that she should spend more time thinking about how best to use Microsoft Word. Sure, there are people who take whole courses in how to change fonts or do mail merge, but anyone moderately clever just opens the thing up and figures out how it works along the way.

This manual began with a growing sense that our own version of this self-taught seat-of-the-

pants approach to computing was hitting its limits. Again and again, we encountered situations like:

- In trying to replicate the estimates from an early draft of a paper, we discover that the code that produced the estimates no longer works because it calls files that have since been moved. When we finally track down the files and get the code running, the results are different from the earlier ones.

- In the middle of a project we realize that the number of observations in one of our regressions is surprisingly low. After much sleuthing, we find that many observations were dropped in a merge because they had missing values for the county identifier we were merging on. When we correct the mistake and include the dropped observations, the results change dramatically.

- A referee suggests changing our sample definition. The code that defines the sample has been copied and pasted throughout our project directory, and making the change requires updating dozens of files. In doing this, we realize that we were actually using different definitions in different places, so some of our results are based on inconsistent samples.

- We are keen to build on work a research assistant did over the summer. We open her directory and discover hundreds of code and data files. Despite the fact that the code is full of long, detailed comments, just figuring out which files to run in which order to reproduce the data and results takes days of work. Updating the code to extend the analysis proves all but impossible. In the end, we give up and rewrite all of the code from scratch.

- We and our two research assistants all write code that refers to a common set of data files stored on a shared drive. Our work is constantly interrupted because changes one of us makes to the data files causes the others' code to break.

At first, we thought of these kinds of problems as more or less inevitable. Any large scale endeavor has a messy underbelly, we figured, and good researchers just keep calm, fight through the frustrations, and make sure the final results are right. But as the projects grew bigger, the problems grew nastier, and our piecemeal efforts at improving matters—writing handbooks and protocols for our RAs, producing larger and larger quantities of comments, notes, and documentation—proved ever more ineffective, we had a growing sense that there must be a way to do better.

In the course of a project involving a really big dataset, we had the chance to work with a computer scientist who had, for many years, taught the course on databases at the University of Chicago. He showed us how we could organize our really big dataset so that it didn't become impossible to work with. Neat, we thought, and went home.

Around that time we were in the middle of assembling a small (but to us, very important) dataset of our own. We spent hours debating details of how to organize the files. A few weeks in we realized something. We were solving the same problem the computer scientist had shown us how to solve. Only we were solving it blind, without the advantage of decades of thought about database design.

Here is a good rule of thumb: If you are trying to solve a problem, and there are multi-billion dollar firms whose entire business model depends on solving the same problem, and there are whole courses at your university devoted to how to solve that problem, you might want to figure out what the experts do and see if you can't learn something from it.

This handbook is about translating insights from experts in code and data into practical terms for empirical social scientists. We are not ourselves software engineers, database managers, or computer scientists, and we don't presume to contribute anything to those disciplines. If this handbook accomplishes something, we hope it will be to help other social scientists realize that there are better ways to work.

Much of the time, when you are solving problems with code and data, you are solving problems that have been solved before, better, and on a larger scale. Recognizing that will let you spend less time wrestling with your RA's messy code, and more time on the research problems that got you interested in the first place.

Chapter 2

Automation

Rules

(A) Automate everything that can be automated.

(B) Write a single script that executes all code from beginning to end.

Let's start with a simple research project. We wish to test the hypothesis that the introduction of television to the US increased sales of potato chips. We receive an Excel file by e-mail with two worksheets: (i) "tv," which contains for each county in the US the year that television was first introduced; and (ii) "chips," which contains total sales of potato chips by county by year from 1940 to 1970. We wish to run a panel regression of log chip sales on a dummy variable for television being available with county and year fixed effects.

Here is one way we might proceed: Open the file in Excel and use "Save As" to save the worksheets as text files. Open up a statistical program like Stata, and issue the appropriate commands to load, reshape, and merge these text files. Define a new variable to hold logged chip sales, and issue the command to run the regression. Open a new MS Word file, copy the output from the results window of the statistical program into a table, write up an exciting discussion of the findings, and save. Submit to a journal.

Just about everybody learns early in graduate school, if not before, that this "interactive" mode of research is bad. They learn that the data building and statistical analysis should be stored in scripts—`.do` files in Stata, `.m` files in Matlab, `.r` files in R, and so forth.

It is worth pausing to remember why we don't like the interactive mode. There are many

reasons, but two big ones.

The first is replicability. If the next day, or the next year, we want to reproduce our regression of chip sales on TV, we might dig up `tv.csv` and `chips.csv`, load them back into Stata, and set to work reshaping, merging, defining variables, and so forth. Perhaps we will get lucky, since this analysis is so simple, and get back the same coefficient when we run the regression. Or perhaps not. Even in this simple example, there are innumerable things that could go wrong: since writing the paper we have received an updated version of `tv.csv` and we inadvertently use the new one rather than the old one; we forget that we dropped several counties whose chip sales were implausibly large; we compute regular standard errors whereas before we computed robust standard errors; and so on.

On a deeper level, because there is no record of the precise steps that were taken, there is no authoritative definition of what the numbers in our paper actually are. If someone later asks why the number of observations reported in our table is different from the number of observations in the raw data, or how we computed our standard errors, or what we did with county-years with missing chip sales, and we ran the analysis interactively, we will have no way to say for sure.

The second reason is efficiency. If we decide to run a different regression, say using the level rather than the log of chip sales, we will have to go back and repeat all of the steps of building and cleaning the data. We can avoid this by saving the combined dataset before running any regressions, but if we later wish to change which observations we keep and which we drop, we will be back to square one.

In a real project, there might be a thousand steps from raw data to final results. For each of these, there could be several alternatives, detours, and experiments that were tried and discarded. Each step is typically run hundreds of times as the analysis is developed and refined. Trying to run and re-run all these steps interactively would be completely untenable.

For this reason, most researchers learn to script key steps, especially data manipulation and statistical analysis. Here is what the project directory for the paper above might look like after we switched to writing .do files, expanded our analysis a bit, and switched to LaTeX for word processing:

```
chips.csv        mergefiles.do         tv_potato_submission.pdf

cleandata.do     regressions_alt.do    tv_potato.tex

extract0B.xls    regressions_alt.log   tv.csv

fig1.eps         regressions.do        tvdata.dta

fig2.eps         regressions.log

figures.do       tables.txt
```

This is certainly a big improvement over our initial interactive approach. If we stare at these files for a while, we can probably work out more or less what they are. `Extract0B.xls` is the raw data file, `chips.csv` and `tv.csv` are the text files exported from Excel, and `tvdata.dta` is the combined data file in Stata format. `Mergefiles.do` and `cleandata.do` are the scripts that build the data, `figures.do`, `regressions.do`, and `regressions_alt.do` are the scripts that run the analysis, and the `.log` and `.eps` files are the output. `Tv_potato.tex` is the paper, `tables.txt` contains the tables, and `tv_potato_submission` is the PDF version we submitted to the journal.

But if we set about actually trying to reproduce `tv_potato_submission.pdf`, we'd immediately run into a bunch of questions. Should we export all observations from `extract0B.xls`, or just those with nonmissing data? Which should be run first, `cleandata.do` or `mergedata.do`? Does it matter in which order we run `regressions.do` and `figures.do`? Is the output from `regression_alt.do` actually used in the paper or is this file just left over from some experimentation? What is `tables.txt`? Is it produced manually or by code? Which numbers in the log files correspond to the numbers reported in the paper? Is `tv_potato_submission.pdf` just a PDF version of `tv_potato.tex` or did we do additional formatting, etc. before submitting to the journal?

We suspect that the experience of trying to reverse-engineer the build steps for a directory like this will feel familiar to many readers who have tried to make sense of directories their RAs or coauthors produced, or even directories that they produced themselves a few months in the past. In this toy example, the problems are probably surmountable and, assuming that we didn't do anything silly like modify and rerun `regressions.do` after the PDF was produced, we could probably reproduce the paper in a reasonable amount of time. But as most of us know from painful experience, the reverse-engineering process for a moderately complex project can easily become days or weeks of frustrating work, and the probability of those "silly" mistakes that render

replication all but impossible is remarkably high.

To make the output of our directory replicable, we need to automate more steps. And we need a way to store the information about the order in which the steps are run.

First, let's add a Stat/Transfer script called `export_to_csv.stc` that handles the conversion from Excel. (Stata can also do this directly using the "import excel" command.) Next, let's switch from outputting `tables.txt` to outputting `tables.tex`, a LaTeX file produced by Stata's "outreg" command.

Finally, let's add another key script to the directory, called `rundirectory.bat`, which is a Windows shell script. Its contents look like this:

```
---- rundirectory.bat ----
stattransfer export_to_csv.stc
statase -b mergefiles.do
statase -b cleandata.do
statase -b regressions.do
statase -b figures.do
pdflatex tv_potato.tex
```

The `rundirectory.bat` script works like a roadmap, telling the operating system how to run the directory. Importantly, the rundirectory script also tells a human reader how the directory works. But unlike a readme file with notes on the steps of the analysis, `rundirectory.bat` cannot be incomplete, ambiguous, or out of date.

The proof in the pudding is that we can now delete all of the output files in the directory – the `.csv` files, the `.log` and `.eps` files, `tables.tex`, the `.pdf` – and reproduce them by running `rundirectory.bat`. This is the precise sense in which the output is now replicable.

Writing a shell script like `rundirectory.bat` is easy.[1] You may need a few tweaks, such as adding Stata to your system path, but many of these will be useful anyway. You could write all these steps into a Stata script (`rundirectory.do`), but a system shell provides a more natural interface for calling commands from multiple software packages, and for operating system commands like moving or renaming files.

---

[1]If you don't use Windows, Linux shell files work almost the same way. And if you're comfortable with Python, you can do even better, and write a `rundirectory.py` that will work on both Windows and Linux systems.

Of course, `rundirectory.bat` does not automate *everything*. We could (and, admittedly, are tempted to) write a little Python script to submit the paper to a journal, but that seems like overkill even to us.

On the other hand, we have consistently found that pushing the boundaries of automation pays big dividends. The costs tend to be lower than they appear, and the benefits bigger. A rule of research is that you will end up running every step more times than you think. And the costs of repeated manual steps quickly accumulate beyond the costs of investing once in a reusable tool.

We used to routinely export files from Excel to CSV by hand. It worked ok until we had a project that required exporting 200 separate text files from an Excel spreadsheet. We followed our usual practice and did the export manually. Some time later, the provider sent us a new Excel file reflecting some updates to the data. We had learned our lesson.

Chapter 3

Version Control

Rules

(A) Store code and data under version control.

(B) Run the whole directory before checking it back in.

In the last chapter, we showed what the project directory for our seminal TV and potato chips project might look like. After we work on the directory for a while, the key files might look like this:

```
cleandata_022113.do     cleandata_022613.do       regressions.log
cleandata_022113a.do    cleandata_022613_jms.do   regressions_022413.do
chips.csv               tvdata.dta                regressions_022713_mg.do
regressions_022413.log
```

Dates are used to demarcate versions of files. Initials (JMS for Jesse, MG for Matt) are used to indicate authorship.

There are good reasons to store multiple versions of the same file. The most obvious is that it provides a quick way to roll back changes you want to discard. Another is that it facilitates comparison. Maybe Matt wants to show Jesse how he's thinking of changing their main specification. Creating `regressions_022713_mg.do` may be a good way to illustrate what he has in mind. If Jesse doesn't like it, he can always delete it.

The goal is admirable, but the method is wrong. There are two main reasons why. First, it is a pain. The researcher needs to decide when to "spawn" a new version and when to continue to edit the old one (hence `022113a`). The researcher needs to tag authorship and date every file. Failing to do that will result in confusion: Why is the date on the file name February 21 when the operating system says this was last edited in March?

And confusion is the second, and by far the more important, reason why this "date and initial" method is poor. Look at the file names above and answer the following questions: Which is the log file produced by `regressions_022713_mg.do`? Did the author (darn you, Matt!) fail to change the output file name in the code, overwriting `regressions_022413.log`? Did he simply not output a log?

Which version of `cleandata.do` produces the data file used by `regressions_022413.do`? Is it the one labeled `022113a`—the last one before February 24? Or was `regressions_022413` created on February 24 but edited later, raising the possibility that it needs output from `cleandata_022613.do` to run correctly? Unfortunately, we failed to tag `tvdata.dta` with a date and initial—probably because changing the file name in three different places with each new version is an enormous hassle.

Given a few minutes to look at the system dates and file contents, you could probably work out which inputs are needed by which scripts. And, having learned your lesson, next time you will harangue your coauthor and RAs to remember to date and initial every script, LOG file, and intermediate data file, so hopefully there's no more confusion.

This is too much work just to keep track of multiple versions of files. And it creates a serious risk that, later, you won't be able to sort out which file goes with which, and hence you won't be able to replicate your results. Fortunately, your computer can take care of this for you, automatically, using free software that you can set up in a few minutes.

Before we tell you how, we will start with a fact. (This is, after all, a handbook for empirical researchers.) *Not one piece of commercial software you have on your PC, your phone, your tablet, your car, or any other modern computing device was written with the "date and initial" method.*

Instead, software engineers use a tool called version control software to track successive versions of a given piece of code. Version control works like this. You set up a "repository" on your PC (or, even better, on a remote server). Every time you want to modify a directory, you "check it

out" of the repository. After you are done changing it, you check it back in. That's it. You don't change file names, add dates, or anything. Instead, the software remembers every version that was ever checked in.

What happens if you change your mind about something? You ask the software for a history of changes to the directory and, if you want to go back to an old version of the directory or even of a single file, the operation just takes a click.

And what about your sneaky coauthor's decision to change the main regression model specification without telling you? The version control software automatically records who authored every change. And if you want to see what the changes were, most modern packages will show you a color-coded side-by-side comparison illustrating which lines of code changed and how.[1]

The main thing about this approach that is great, and the reason real software engineers *must* use a tool like this, is that it maintains a single, authoritative version of the directory at all times. In rare cases where two people try to make simultaneous and conflicting changes to the same file, the software will warn them and help them reconcile the conflicts.

A major ancillary benefit is therefore that you can edit without fear. If you make a mistake, or if you start in a new direction but later change your mind, you can always roll back all or part of your changes with ease. This requires no keeping track of dates and initials. All file names can remain just as nature intended. The software handles the versioning for you, so you can focus on writing the code and making it right. You didn't spend six years in grad school so you could type in today's date all over the place.

To visualize how much better your life would be with your code and data under version control, recall (if, gasp, you are old enough) what word processing was like before the invention of the "undo" command. A bad keystroke might spell doom. Version control is like an undo command for *everything*.

So our first rule is to keep everything—code and data—under version control. In fact, version control is fantastic for things like drafts of your papers, too. It allows you to overwrite changes without fear, to keep track of authorship, etc. (Some readers will have noticed the attractions of the "version history" feature of GoogleDocs, which is based on version control models from software

---

[1]For a slightly more advanced user, there are also well-defined methods for changing code in a way that is explicitly tentative, so you can "pencil in" some changes and let your coauthor have a look before you take them on board.

engineering. With version control software you get that functionality with LATEX, LYX, or whatever is your favorite editing package.)

But if you want to get the most out of this approach, there is a second rule: you have to run the entire directory before you check in your changes. Return to our example, now dropping the annoying date and initial tags, and adding back `rundirectory.bat` (which, you'll recall, will run every script from top to bottom).

```
rundirectory.bat   tvdata.dta
cleandata.do       regressions.do
chips.csv          regressions.log
```

Suppose Jesse modifies `cleandata.do` and runs it to overwrite `tvdata.dta`. If Jesse checks in that change, Matt may find later that `regressions.do` breaks when he tries to run it, because of a change to `tvdata.dta` that `regressions.do` wasn't expecting.

The way to fix this problem and ensure it never happens again is just to execute `rundirectory.bat`, from start to finish, and check for errors before checking in the directory. If every version you check in has been run successfully via `rundirectory.bat`, then you know that, barring changes in the software itself, the next time you check it out, you will get back the output in `regressions.log` exactly.

Note that this is not a problem with the version control software. The "date and initial" method creates the same potential for this type of within-directory conflict, arguably more so, since a lot of effort is required to keep track of which input files are required for which scripts. Rather, version control, coupled with the rule of checking in complete runs of a directory, provides a comprehensive solution that guarantees both replicability and undo-ability with minimal effort.

OK, you're convinced. Now what? A step-by-step guide to setting up and using version control software is a bit outside our scope here. But for what it's worth, we use a SubVersion repository that we interact with using the very nice TortoiseSVN browser for Windows. Comparable software exists for Macs. More recent version control methods like Git or BitBucket may be worth checking out. It will probably take you a couple days to set up a repository and learn how you want to interact with it. You will break even on that time investment within a month or two.

Chapter 4

Directories

Rules

(A) Separate directories by function.

(B) Separate files into inputs and outputs.

(C) Make directories portable.

Let's return to the main directory for our potato chip project:

```
                 ---C:/tv_and_potato/---
chips.csv       mergefiles.do        tv_potato_submission.pdf
cleandata.do    regressions_alt.do   tv_potato.tex
extractOB.xls   regressions_alt.log  tv.csv
fig1.eps        regressions.do       tvdata.dta
fig2.eps        regressions.log      rundirectory.bat
figures.do      tables.txt           export_to_csv.stc
```

The directory above contains all the steps for the entire project, governed by the single batch file `rundirectory.bat`.

Having a single directory that has and does everything has some appeal, but for most real-world research projects this organizational system is not ideal. Consider the following scenarios. (i) The researcher wants to change a regression specification, but does not want to re-run the entire data build. (ii) The researcher learns about a neat Stata command that makes the script

`export_to_csv.stc` and its outputs `tv.csv` and `chips.csv` unnecessary. Before making this improvement, however, the researcher must search through `regressions.do` and `regressions_alt.do` to make sure these scripts do not depend on `tv.csv` and `chips.csv` in addition to `tvdata.dta`.

Consider the following alternative directory structure (leaving aside the TEX and PDF files for simplicity):

```
---C:/build---              ---C:/analysis---
/input                      /input
    extract0B.xls               tvdata.dta (link to C:/build/output)


/code                       /code
    rundirectory.bat            rundirectory.bat
    export_to_csv.stc           regressions.do
    mergefiles.do               regressions_alt.do




/output                     /output
    tvdata.dta                  fig1.eps
                                fig2.eps
                                tables.txt



/temp                       /temp
    chips.csv                   regressions.log
    tv.csv                      regressions_alt.log
```

There are now two high-level directories. One contains the code to build a useable Stata file from the raw inputs. The other contains code to take the Stata file and turn it into figures and tables for the paper.

Within each high-level directory there is a consistent subdirectory structure that separates inputs, outputs, code, and temporary or intermediate files. Each directory is still controlled by a single script called `rundirectory.bat` that executes from start to finish. (In fact, we advocate

having `rundirectory.bat` start by clearing the contents of */temp* and */output*, so you can be sure all your output is produced by your current code.)

It is now easy to modify the analysis without re-running the data build. And it is now clear from the directory structure that only `tvdata.dta` is required by the analysis code: `chips.csv` and `tv.csv` are explicitly designated as temp files. Finally, because we are using a local link to the input data, we can write all the code in the analysis directory to use local references (`../input/tvdata.dta` instead of `C:/build/output/tvdata.dta`).[1]

A disadvantage of this structure is that the code in C:/analysis will break if it is run on a different machine (where the link to tvdata.dta is not valid), or if the structure of `tvdata.dta` changes. What we actually do in practice is therefore slightly more complicated: in place of `tvdata.dta`, we link to *fixed revisions* of our datasets on shared network storage. This means an /analysis/ directory like the one above can be run anywhere with network access. The fixed revision bit means that if someone modifies the structure of the data and checks in a new revision, the analysis code will continue to work, as it still points to the old revision. (Of course at some point someone will probably want to redirect it to the new revision, but the user gets to decide when to do this, rather than having her code break unexpectedly.)

We have only outlined a few of the advantages of using modular, functional directories to organize code. There are many others. For example, the output of *C:/build* is now easily accessible by any directory, which makes it easier to have multiple projects that use the same data file without creating multiple, redundant copies. And, separating scripts into functional groups makes debugging easier and faster when something goes wrong.

---

[1]It is easy to create a "symbolic link" to a file in another directory; see your operating system's documentation for details. An alternative would be to add code to `rundirectory.bat` that copies `tvdata.dta` into `/input` from `C:/build/output.` This would also allow your code to use local file references, but at the cost of duplicating the storage of `tvdata.dta.`

Chapter 5

Keys

Rules

(A) Store cleaned data in tables with unique, non-missing keys.

(B) Keep data normalized as far into your code pipeline as you can.

It is well known that television went to big cities first. So a good analysis of the effect of television on potato chip consumption requires good data on population as a control. We ask an RA to prepare a population dataset to facilitate our analysis. Here it is:

county	state	cnty_pop	state_pop	region
36037	NY	3817735	43320903	1
36038	NY	422999	43320903	1
36039	NY	324920	.	1
36040	.	143432	43320903	1
.	NY	.	43320903	1
37001	VA	3228290	7173000	3
37002	VA	449499	7173000	3
37003	VA	383888	7173000	4
37004	VA	483829	7173000	3

What a mess. How can the population of the state of New York be 43 million for one county

but "missing" for another? If this is a dataset of counties, what does it mean when the "county" field is missing? If region is something like Census region, how can two counties in the same state be in different regions? And why is it that all the counties whose codes start with 36 are in New York except for one, where the state is unknown?

We can't use these data, because we don't understand what they mean. Without looking back at the underlying code, we could never say confidently what every variable is, or even every row. And we can forget trying to merge on attributes from another dataset. How would we know which state goes with county 36040? Or which region to use for 37003?

We know many researchers who spend time wrestling with datasets like this, and barking at RAs, students, or collaborators to fix them.

There must be a better way, because we know that large organizations like financial institutions, retailers, and insurers have to manage much more complex data in real time, with huge consequences of mistakes.

Long ago, smart people figured out a fundamental principle of database design: that the physical structure of a database should communicate its logical structure.

If you gave your county Census data to someone with training in databases, you'd probably get back something like this, called a *relational database*:

county	state	population
36037	NY	3817735
36038	NY	422999
36039	NY	324920
36040	NY	143432
37001	VA	3228290
37002	VA	449499
37003	VA	383888
37004	VA	483829

state	population	region
NY	43320903	1
VA	7173000	3

Now the ambiguity is gone. Every county has a population and a state. Every state has a population and a region. There are no missing states, no missing counties, and no conflicting definitions. The database is self-documenting. In fact, the database is now so clear that we can

forget about names like county_pop and state_pop and just stick to "population." Anyone would know which entity's population you mean.

Note that when we say relational database here, we are referring to how the data are structured, not to the use of any fancy software. The data above could be stored as two tab-delimited text files, or two Stata `.dta` files, or two files in any standard statistical package that expects rectangular data.

Stepping back from this example, there are a few key principles at work here. To understand these it is helpful to have some vocabulary. Data is stored in rectangular arrays called *tables*. In the example above, there is a county table and a state table. We will refer to rows of tables as *elements* and columns of tables as *variables.*

Critically, each table has a *key* (rule A). A key is a variable or set of variables that uniquely identifies the elements of a table. The variables that form the key never take on missing values, and a key's value is never duplicated across rows of the table. So, a state table has one and only one row for New York, and no rows where state is missing.

Each variable in a table is an attribute of the table's elements. County population is a property of a county, so it lives in the county table. State population is a property of a state, so it cannot live in the county table. If we had panel data on counties, we would need separate tables for things that vary at the county level (like state) and things that vary at the county-year level (like population). This is the main reason that most data files researchers use for analysis do not satisfy the database rules: they typically combine variables defined at several different levels of aggregation.

The variables in a table can include one or more *foreign keys*. A foreign key is the key for another table in the database. In the county table, the variable "state" is a foreign key: it matches a county to an element of the state table. Foreign keys obey the same rules as all variables: they are included at the level of logical attribution. States are in regions, and counties are in states, so "state" shows up in the county table, and "region" in the state table.

Data stored in the form we have outlined is considered *normalized.* Storing normalized data means your data will be easier to understand and it will be harder to make costly mistakes.

Most statistical software won't run a regression on a relational database. To perform our analysis we are going to need to merge (or join, in database-speak) the tables together to produce a single rectangular array. Plus, we might need to calculate some variables that aren't in our source

data, such as the log of population.

To get from the data you downloaded, entered, or bought from an original source to the matrix on which you will perform estimation, we recommend proceeding in three steps.

First, store your raw data in normalized files that preserve the information in the original data source and follows the rules above. Don't worry about how you plan to use the data. Rather, imagine that you are preparing the data for release to a broad group of users with differing needs. Do this because you, yourself, are likely to want to use the data in ways you do not currently anticipate.

Second, construct a second set of normalized files that includes the transformations of the original variables that you will need for your analysis. For example, you might add to the county table a variable indicating the county's population rank within its state. At this stage you can also bring in variables from other databases. For example, you might use a geography database to bring in county latitude and longitude.

Third, merge together the tables in the database to form the rectangular array on which you will estimate your model. At this stage, your database should still have unique, non-missing keys, but it will likely not be normalized. In our example, you will have a county-level file that includes variables like region that are not properties of counties. If you had panel data, your file would include both county-level and county-year-level variables. Do no data manipulation in this step. If your analysis requires the log of state population, calculate it while your database complies with the rules.

Following the steps above means you can keep your data in a normalized form until the last possible step in the code (rule B).

Chapter 6

Abstraction

Rules

(A) Abstract to eliminate redundancy.

(B) Abstract to improve clarity.

(C) Otherwise, don't abstract.

We are concerned about spatial correlation in potato chip consumption. We want to test whether per capita potato chip consumption in a county is correlated with the average per capita potato chip consumption among other counties in the same state. First we must define the "leave-out" mean of per capita consumption for each county:

```
egen total_pc_potato = total(pc_potato), by(state)
egen total_obs = count(pc_potato), by(state)
gen leaveout_state_pc_potato = (total_pc_potato - pc_potato) / (total_obs - 1)
```

We can now test whether `pc_potato` is correlated with `leaveout_state_pc_potato`. If so, we may need to adjust how we compute the standard errors in our model. We perform our analysis and are comforted to find little evidence of spatial correlation.

But what if we are using the wrong level of aggregation? Maybe spatial correlation will show up at the level of the metropolitan area. Let's copy and paste the code above and then adapt it to use metropolitan area instead of state as the level of aggregation:

```
egen total_pc_potato = total(pc_potato), by(metroarea)
egen total_obs = count(pc_potato), by(state)
gen leaveout_metro_pc_potato = (total_pc_potato - pc_potato) / (total_obs - 1)
```

And while we're at it, let's check if there is more spatial correlation in potato chip consumption when measured on a per-household rather than per-capita basis. For this we will need a third leave-out mean:

```
egen total_hh_potato = total(hh_potato), by(metroarea)
egen total_obs = count(hh_potato), by(state)
gen leaveout_metro_hh_potato = (total_hh_potato - pc_potato) / (total_obs - 1)
```

Note the errors. In the first "copy-and-paste" operation, we failed to replace an instance of `state` with `metroarea`. In the second, we propagated the first error, plus we failed to replace one use of the per-capita potato variable with the per-household analogue. The code will run, but everything after the first code block will be totally wrong.

Consider an alternative to the copy-and-paste approach, which is to write a general-purpose function that computes the leave-out mean of a variable:

```
program leaveout_mean

    syntax, invar(varname) outvar(name) byvar(varname)
    tempvar tot_invar count_invar
    egen 'tot_invar'= total('invar'), by('byvar')
    egen 'count_invar'= count('invar'), by('byvar')
    gen 'outvar' = ('tot_invar' - 'invar') / ('count_invar' - 1)

end
```

Having defined the function above, we can now replace our three code blocks with three lines:

```
leaveout_mean, invar(pc_potato) outvar(leaveout_state_pc_potato) byvar(state)
leaveout_mean, invar(pc_potato) outvar(leaveout_metro_pc_potato) byvar(metro)
leaveout_mean, invar(hh_potato) outvar(leaveout_metro_hh_potato) byvar(metro)
```

Now the amount of copying and pasting is minimized: each input is changed only once as we go from line to line. And because we wrote the `leaveout_mean` function to be totally general, we can use it for other projects as well as this one.[1] We will never again have to write code from scratch to compute a leave-out mean.

Key to achieving these goals is recognizing that all three code blocks were just specific instances of the same abstract idea: compute the mean of a variable across observations in a group, excepting the given observation. In programming, turning the specific instances of something into a general-purpose tool is known as *abstraction*.

Abstraction is essential to writing good code for at least two reasons. First, as we saw above, it eliminates redundancy, which reduces the scope for error and increases the value you can get from the code you write. Second, just as importantly, it makes code more readable.[2] A reader scanning one of the three code blocks above might easily miss their purpose. By contrast, a call to a function called `leaveout_mean` is hard to misunderstand.

Abstraction can be taken too far. If an operation only needs to be performed once, and the code that performs it is easy to read, we would not advise abstraction. Abstracting without a purpose can lead you to spend a lot of time dealing with cases that will never come up in your work.

When you do have a function you plan to use often, you should take the time to implement it carefully. One thing we have found helpful is the software engineering practice of "unit testing." This means writing a script that tests out the behavior of the function you've written to make sure it works as intended. For example, we might make some fake data and verify that the `leaveout_mean` calculates the leave-out mean correctly. An advantage of unit testing is that it allows you to safely change your function without fear that you will introduce errors that will break your code down the line. It also provides a convenient way to document how the function works: what inputs it requires, what inputs it will not accept, etc.

Abstraction is not just about code. It is relevant anywhere you find yourself repeating an operation. The principles in this chapter, for example, explain why word processing packages come with templates for standard document types like memos or reports. And these principles are

---

[1]In Stata, as with just about any program you are likely to use, it is easy to make a function portable and accessible anytime you use the program.

[2]In fact, we have found that the general version of a function is often easier to write as well as easier to read. (To see why, think about how much harder it would be to program a linear regression for a specific matrix of variables than for a general one.)

the reason why, rather than just repeatedly telling our RAs how we thought code should be written, we decided to write this handbook!

Chapter 7

Documentation

Rules

(A) Don't write documentation you will not maintain.

(B) Code should be self-documenting.

We have estimated the effect of television on potato chip consumption. To illustrate the pernicious consequences for society we wish to perform a welfare analysis, for which we will need to compute an elasticity. Fortunately, Jesse's dissertation studied the effect of a tax increase on demand for potatoes, from which we can back out the elasticity of demand.

Here is how a section of our Stata code might look:

```
* Elasticity = Percent Change in Quantity / Percent Change in Price
* Elasticity = 0.4 / 0.2 = 2
* See Shapiro (2005), The Economics of Potato Chips,
* Harvard University Mimeo, Table 2A.
compute_welfare_loss, elasticity(2)
```

Notice the helpful comments that provide a roadmap to the reader.

Many researchers we know spend a lot of energy haranguing themselves, their coauthors, and their research assistants to write more comments like the above, and in general, to carefully docu-

ment the organization of their code and data outside of the scripts and data files themselves. You might expect us to say the same. After all, we love organizing things. But in this chapter we will try to convince you to document less, not more. To see why, we continue our story.

A few months after writing the first version of our script we return to our code to revise the analysis. We find the following:

```
* Elasticity = Percent Change in Quantity / Percent Change in Price
* Elasticity = 0.4 / 0.2 = 2
* See Shapiro (2005), The Economics of Potato Chips,
* Harvard University Mimeo, Table 2A.
compute_welfare_loss, elasticity(3)
```

Notice the conflict in red. Maybe someone noticed a typo in the original calculation, or decided to use the estimates from Table 2B instead of Table 2A of Jesse's dissertation. Whatever its origin, the problem is clear: the comments contradict the code, and it is now unclear which (if either) is correct. Someone will have to go back to the source to figure out what number we should be using.

Readers will notice that this is an instance of a more general problem: anytime you have more than one representation of the same information (in this case, an elasticity), you run the risk that the two will someday come in conflict. In the best case scenario, you will need to do some work to untangle the mess. In the worst case scenario, your results will be wrong or internally inconsistent.

The problem of internal inconsistency is especially severe when it comes to documentation—comments, notes, readmes, etc.—because you don't *have* to keep them up to date for your code to work or for your results to be quantitatively right. It is therefore tempting to make improvements to the code without making parallel improvements to the comments, only to find later that your comments are confusing or misleading. In the case above, the practice of letting comments go stale resulted in code that is probably less clear than it would have been if we had not had so much documentation in the first place.

To avoid such confusion, you will need to keep your comments up to date, meaning just as up to date as your code. If it's not worth maintaining a piece of documentation up to that standard, it probably isn't worth writing it in the first place (rule A).

That raises the important question of how to make the code clear without extensive comments. Imagine the selection above with no comments at all. How would a reader know why the elasticity is 2 and not 3?

To solve that problem we turn to the code itself. Much of the content of the comments above can be readily incorporated into the code:

```
* See Shapiro (2005), The Economics of Potato Chips,
* Harvard University Mimeo, Table 2A.
local percent_change_in_quantity = -0.4
local percent_change_in_price = 0.2
local elasticity = `percent_change_in_quantity'/`percent_change_in_price'
compute_welfare_loss, elasticity(`elasticity')
```

This code block contains just as much documentation as the one we started with. It makes clear both the formula for the price elasticity and the quantitative components we are using. But it is far better than the original code, because it has far less scope for internal inconsistency. You can't change the percent change in quantity without also changing the elasticity, and you can't get a different elasticity number with these percent changes.

When possible, then, you should write your code to be self-documenting (rule B). Use the naming of variables and the structure of the code to help guide a reader through your operations. That's a good idea anyway, because even the best comments can't untangle a coding mess. To boot, writing such code will mean you don't have to write comments and other notes only to find that they have later lost their grip on what the code is really doing.

These principles apply far beyond code, and indeed they underlie many of the other chapters in this handbook. Organizing your data files so that their structure makes their meaning clear lets you avoid pairing every dataset you make with extensive documentation (chapter 5). Naming files, directories, and other objects intelligently means their names declare their function (chapter 4). A cleverly drawn figure or table will often say so much that notes are present only to confirm the obvious or clarify minor details. And so on.

Documentation does have its place. In the example above, if we don't include the citation to Jesse's stellar thesis, how will a reader know where 0.4 comes from? There is no (practical) way

to script the link back to the original paper, so a comment is appropriate.

Documentation can be used to make clear that something is right when it at first may seem wrong. Suppose, for example, we have a variable *y* distributed lognormal with location μ and scale σ. If we wish to compute the log of the variable's expectation, it might be wise to write

```
* Log of the expectation is not the expectation of the log
* See http://en.wikipedia.org/wiki/Log-normal_distribution
log_expected_y = 'mu' + 0.5*('sigma'^2)
```

so that a reader isn't surprised that the expression is not simply $\log(E(y)) = \mu$. Of course, what to document is in the eye of the beholder: if you and your collaborators are not likely to forget the expression for the expectation of a lognormal, then the comment above is probably superfluous.

Documentation can also be used to prevent unintended behavior. Suppose you write a command to estimate a regression model via maximum likelihood. If two or more variables are collinear, your solver will iterate forever. So, you may wish to put a warning in the code: "Don't try to estimate an unidentified model." But be careful. As we note above, nothing documents code quite like code. Writing a function to test whether your $(X'X)$ matrix has full rank will provide just as much documentation, will not require the user to be conscientious enough to read the comments, and will likely lead to a faster resolution of the problem.
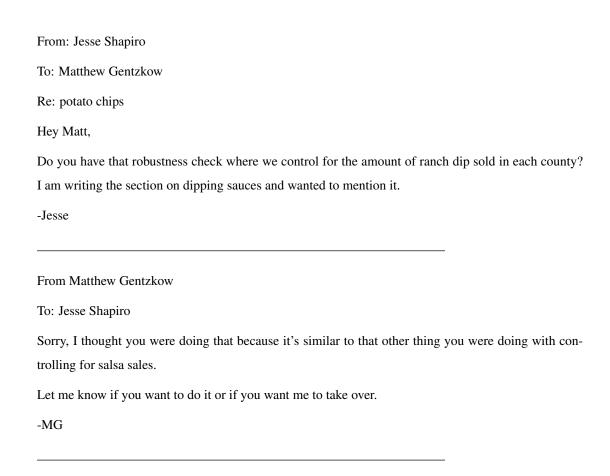
Which brings us to a related point. In Jesse's house there is a furnace room with two switches. One controls a light. The other turns off the hot water for the whole house. When he first moved in, people (let's not name names) conducting innocent business would occasionally shut off the hot water while fumbling for the light switch. He tried having a sign: "Do not touch this switch." But in the dark, in a hurry, a sign is worthless. So he put a piece of tape over the switch. If there are some inputs you really, really want to prevent, comments that say "don't ever do X" are not the way to go. Write your code so it will not let those inputs in the door in the first place.

Chapter 8

Management

Rules

(A) Manage tasks with a task management system.

(B) E-mail is not a task management system.

From: Jesse Shapiro

To: Matthew Gentzkow

Re: potato chips

Hey Matt,

Do you have that robustness check where we control for the amount of ranch dip sold in each county? I am writing the section on dipping sauces and wanted to mention it.

-Jesse

From Matthew Gentzkow

To: Jesse Shapiro

Sorry, I thought you were doing that because it's similar to that other thing you were doing with controlling for salsa sales.

Let me know if you want to do it or if you want me to take over.

-MG

From: Jesse Shapiro

To: Matthew Gentzkow, Michael Sinkinson

I thought Matt was doing ranch dip and Mike was doing salsa?

-Jesse

From: Michael Sinkinson

To: Matthew Gentzkow, Jesse Shapiro

I did the salsa robustness check two weeks ago. See my e-mail from 8/14, 9:36am.

-Mike

From: Jesse Shapiro

To: Michael Sinkinson, Matthew Gentzkow

Right, but in that e-mail you were controlling for the log of salsa consumption. I thought we agreed we wanted the level of consumption?

-Jesse

From: Michael Sinkinson

To: Jesse Shapiro, Matthew Gentzkow

On it!

-Mike

What's wrong with this picture? Mainly, it's ambiguity. Mike thought his task was done, when Jesse thought it was not. Matt thought Jesse was working on the ranch dressing task, but Jesse thought Matt was doing it. In fact, a careful reader will notice that even after all that e-mail, it's *still* not clear who is going to do the ranch dressing robustness check!

It's worse than that. If we come back to the salsa task in two weeks, where will we look to find out its status? This thread? The one Mike mentions from 8/14? And how will we reference our discussion? By date? By forwarding this whole thread, including all the extraneous exchanges about ranch dressing?

If you work alone, these problems are small. You probably have a legal pad or a Word document or a spot on your whiteboard where you keep track of what you need to do. Every now and again you might forget what you were planning to do or where you jotted something down, but if you are organized you probably get by ok.

The minute two people need to work together, however, the problems exemplified in the thread above are big. And although we haven't proved this formally, we think they grow more than arithmetically with the number of people (coauthors, RAs, etc.) involved in a project.

Software firms handle project and task management systematically. Microsoft does not just say, "Hey Matt, when you get a chance, can you add in-line spell-checking to Word?"

Rather, enterprises engaged in collaborative work use project and task management systems that enforce organized communication and reporting about tasks. In the old days, those often involved handing physical reports up the chain of command. Now, they increasingly involve the use of browser-based task-management portals.

In one of these portals, Mike's salsa task would have looked like this:

Task: Salsa Robustness Check

Assigned To: Michael Sinkinson

Assigned By: Jesse Shapiro

Subscribed to Comments: Matthew Gentzkow

Status: Completed.

Description:

Run main specifications adding a control for per capita salsa consumption.

Add a line to our robustness table reflecting the results.

Comment by: Michael Sinkinson

On it!

Comment by: Michael Sinkinson

See the new version of the paper posted in /drafts/Potato Chips and the supporting code in /analysis/Potato Chips. Is this what you had in mind?

Comment by: Jesse Shapiro

Almost. Our econometric model implies that salsa consumption should enter in levels not logs. Can you revise?

Comment by: Michael Sinkinson

Ok, how about now?

Comment by: Jesse Shapiro

Yup, looks good.

Completed By: Michael Sinkinson

Notice that now there is no ambiguity about whose responsibility the task is or what the goals are. Anyone looking at the task header will know that Mike is expected to do it, and no explicit communication is needed to figure out who is doing what.

There is also a natural place to store communication about the task. Everyone expects that questions and answers will be posted to the appropriate task. And, weeks, months, or years later, there will still be a task-specific record of who did what and why.

There are lots of good online systems for task management available at the moment that look something like the example above. Many are free or at least have a free no-frills option. Most have apps for mobile devices and offer some kind of e-mail integration so, for example, Mike's comments above would be e-mailed to Jesse so he knows there's something he needs to look at.

These systems are changing all the time and which one you want is a matter of taste, style, budget, and the like, so we won't review them all here. Good free options as of this writing include Asana (www.asana.com), Wrike (www.wrike.com) and Flow (www.getflow.com). We use a program called JIRA, which is not free and requires a little more work to install.

While we're on the subject of useful tools, you should probably get yourself set up with some kind of collaborative note-taking environment. That way, you're not bound by the limitations of your task management system in what you can share or record. It's helpful to have a place to

jot down thoughts or display results that are less structured than the code that produces your final paper, but more permanent than an e-mail or conversation.

The best system is one that lets you easily organize notes by project and share them with other users. It's great if you can add rich attachments so you can show your collaborators a graph, a code snippet, a table, etc.

There are a bunch of options, and again, many are free. Evernote (www.evernote.com) has a free basic option and is available across lots of platforms and interfaces. Another option for Windows users is OneNote, which is included with Microsoft Office.

Appendix: Code Style

Every piece of code you write has multiple audiences. The most important audience is the computer: if the code does not deliver unambiguous and correct instructions, the result will not be what the author intends.

But code has other audiences as well. Somewhere down the line, you, your coauthor, your RA, or someone wishing to replicate your findings will need to look at the code in order to understand or modify it.

Good code is written with all of these audiences in mind. Below, we collect some of the most important principles that we have learned about writing good code, mostly with examples using Stata or Matlab syntax. A lot has been written about good code style and we don't intend this as a replacement for a good book or more formal training or industry experience. But we have found the notes below useful in reminding ourselves and our collaborators of some important elements of best practice.

Keep it short and purposeful.

No line of code should be more than 100 or so characters long. Long scripts should be factored into smaller functions. Individual functions should not normally be more than 80 or so lines long.

Scripts should not normally be longer than a few hundred lines. If you are finding it hard to make a long script short and purposeful, this is a sign that you need to step back and think about the logical structure of the directory as a whole.

Every script and function should have a clear, intuitive purpose.

Make your functions shy.

A reader should know exactly which variables a function uses as inputs and which variables it can potentially change.

Most functions should explicitly declare their inputs and outputs and should only operate on local variables. Make the set of inputs and outputs as small as possible; the functions should be reluctant to touch any more data than they need to. For example, if a function only depends on the parameter beta, pass it only beta and not the entire parameter vector.

Use global variables rarely if ever.

Order your functions for linear reading.

A reader should be able to read your code from top to bottom without skipping around. Subfunctions should therefore appear immediately after the higher level functions that call them.

Us descriptive names.

Good names replace comments and make code self-documenting.

By default, names for variables, functions, files, etc. should consist of complete words. Only use abbreviations where you are confident that a reader not familiar with your code would understand them and that there is no ambiguity. Most economists would understand that "`income_percap`" means income per capita, so there is no need to write out `income_percapita`. But `income_pc` could mean a lot of different things depending on the context. Abbreviations like `st`, `cnty`, and `hhld` are fine if they are used consistently throughout a body of code. But using `blk_income` to represent the income in a census block could be confusing.

Avoid having multiple objects whose names do not make clear how they are different: e.g., scripts called "state_level_analysis.do" and "state_level_analysisb.do" or variables called `x` and `xx`.

Names can be shorter or more abbreviated when the objects they represent are used frequently and/or very close to where they are defined. E.g., it is sometimes useful to define short names to use in algebraic calculations. This is hard to read:

```
log_coefficient = log((income_percap' * income_percap)^(-1) *///
```

```
income_percap' * log_wage)
```

This is better:

```
X = income_percap
Y = log_wage
log_coefficient = log((X'*X)^(-1)*X'*Y)
```

Pay special attention to coding algebra.

Make sure that key calculations are clearly set off from the rest of the code. If you have a function called `demand()` with 15 lines of setup code, 1 line that actually computes the demand function, and 5 more lines of other code, that 1 line should be set off from the rest of the code or isolated inside a sub-function so it is obvious to a reader scanning the document.

Break complicated algebraic calculations into pieces. Programming languages have no objection to definitions like

```
gen percap_gdp_real = ///
   (consumption + govt_expenditures + exports - imports - taxes) * ///
   10^6 / (price_index * pop_thousands * 1000)
```

or far longer ones. But a human may find it easier to parse the following:

```
gen gdp_millions_nominal = ///
   (consumption + govt_expenditures + exports - imports - taxes)
gen gdp_total_real = gdp_millions * 10^6 / price_index
gen pop_total = pop_thousands * 10^3
gen gdp_percap_real = gdp_total_real / pop_total
```

Complex calculations are better represented in mathematical notation than in code. This is a case where storing documentation (a LaTex or pdf with the calculations written out) alongside code can make sense.

Make logical switches intuitive.

When coding switches, make sure that the conditions are intuitive. Often there is more than one way to express a logical condition. Choosing the most intuitive expression makes the logical meaning of the switch clear, and helps users parse the code quickly. In the following example using Matlab code, suppose x is a vector of 0s and 1s:

```
if max(x) == 0

    y = 0

end
```

This block of code is logically equivalent to:

```
if all(x == 0)

    y = 0

end
```

Both switches check whether the vector x contains only 0s. However, the first condition parses as "if the maximum entry of x if equal to 0", while the second parses as "if all entries of x are zero". The second test is better because it is logically identical to what we want to check, whereas the first test relies on the fact that all entries of x are greater than or equal to 0.

Be consistent.

There are many points of coding style that are mostly a matter of taste. E.g., sometimes people write variable names like `hhld_annual_income` and other times like `hhldAnnualIncome`. Although some people have strong feelings about which is better, we don't. What is important is that everyone on a team use consistent conventions. This is especially important within scripts: if you are editing a program in which all scripts use two-space indent you should use two-space indent too, even if that breaks the normal rule. (Or, you should use grep to update the script to four-space indent).

Check for errors.

Programming languages typically come with debugging tools and informative error handling. Often they are enough, and we do not need to write additional error checking ourselves. For example if in Stata I write

```
gen str x = "hello"
gen y = x^2
```

then Stata will return:

```
type mismatch
r(109);
```

Typically, this will be enough information to alert the user to the fact that the code failed because the user attempted to square a string. Adding additional error-checking to check that x is not a string will add one or more lines of code with little gain in functionality.

However, there are some circumstances in which error-checking should be added to code.

Error-checking should be added for robustness. For example, if the wrong argument will cause your script to become stuck in an infinite loop, or call so much memory that it crashes your computer, or erase your hard drive, you should include code to ensure that the arguments satisfy sufficient conditions so that those outcomes will not occur.

Error-checking should be added to avoid unintentional behavior. For example, suppose function `multiplybytwo()` multiplies a number by 2 but is only written to handle positive reals. For negative reals it produces an incoherent value. Because a user might expect the function to work on any real, it would be a good idea to throw an error if the user supplies a negative real argument. (Of course, it would be even better not to write a function with such confusing behavior.)

Error-checking should be added to improve clarity of error messages. For example, suppose that function `norm()` requires a vector input. Suppose the default error message that gets returned in the event you pass norm() it a scalar is *"Input to '*' cannot be an empty array."* A user could spend a lot of time trying to understand the source of the error. If you suspect that people may be tempted to pass the function a scalar, it is probably worth checking the input and returning a more

informative error message like "*Input to norm() must be a vector.*" Chances are that this will save a few more 30-minute debugging sessions in the future and be worth the time.

Note that there is an intrinsic tradeoff between time spent coding error-handling and time spent debugging. It is not efficient to code explicit handling of all conceivable errors. For example, it is probably not worth adding a special warning in the case where the user passes a string to `norm()`, because the user is unlikely to make that mistake in the first place.

Error checking code should be written so it is easy to read. It should be clearly separated from other code, either in a block at the top of a script or in a separate function. It should be automated whenever possible. If you find yourself writing a comment of the form

```
% Note that x must be a vector
```

ask yourself whether you can replace this with code that throws an error when `isvector(x)` is false. Code is more precise than comments, and it lets the language do the work.

The usual warning against redundancy applies to error-checking. If you have a large program many of whose functions operate on a data matrix X, and there are various conditions that the data matrix must satisfy, write a function called `is_valid_data_matrix()` rather than repeating all the validation checks at the top of each function.

Write tests.

Real programmers write "unit tests" for just about every piece of code they write. These scripts check that the piece of code does everything it is expected to do. For a demand function that returns quantity given price, for example, the unit test might confirm that several specific prices return the expected values, that the demand curve slopes down, and that the function properly handles zero, negative, or very large prices. For a program to compile, it must pass all the unit tests. Many bugs are thus caught automatically. A large program will often have as much testing code as program code.

Many people advocate writing unit tests *before* writing the associated program code.

Economists typically do not write unit tests, but they test their code anyway. They just do it manually. An economist who wrote a demand function would give it several trial values interactively to make sure it performed as expected. This is inefficient because writing the test would

take no more time than testing manually, and it would eliminate the need to repeat the manual tests every time the code is updated. This is just a special case of the more general principle that any manual step that can be turned into code should be.

We therefore advocate writing unit tests wherever possible.

Profile slow code relentlessly.

Languages like Matlab and R provide sophisticated profiling tools. For any script for which computation time is an issue (typically, anything that takes more than a minute or so to run), you should profile frequently. The profiler often reveals simple changes that can dramatically increase the speed of the code.

Profiling code in Stata or similar statistical packages is more difficult. Often, the sequential nature of the code means it is easy to see where it is spending time. When this is not the case, insert timing functions into the code to clarify which steps are slow.

Speed can occasionally be a valid justification for violating the other coding principles we articulate above. Sometimes we are calling a function so many times that the tiny overhead cost of good, readable code structure imposes a big burden in terms of run-time. But these exceptions are rare and occur only in cases where computational costs are significant.

Store "too much" output from slow code.

There is an intrinsic tradeoff between storage and CPU time. We could store no intermediate data or results, and rerun the entire code pipeline for a project back to the raw data stage each time we change one of the tables. This would save space but would require a tremendous amount of computation time. Or, we could break up a project's code into hundreds of directories, each of which does one small thing, and store all the intermediate output along the way. This would let us make changes with little computation time but would use a lot of storage (and would likely make the pipeline harder to follow). Usually, we compromise, aggregating code into directories for conceptual reasons and to efficiently manage storage.

It is important to keep this tradeoff in mind when writing slow code. Within reason, err on the side of storing too much output when code takes a long time to run. For example, suppose you write a directory to estimate several specifications of a model. Estimation takes one hour. At the time

you write the directory, you expect to need only one parameter from the model. Outputting only that parameter is a mistake. It will (likely) be trivial to store estimates of all the model parameters, and the benefits will be large in compute time if, later, you decide it would be better to report results on two or three of the model's parameters.

If estimation is instantaneous, re-estimating the model later to change output format will not be costly in terms of compute time. In such cases, concerns about clarity and conceptual boundaries of directories should take priority over concerns about CPU time.

Separate slow code from fast code.

Slow code that one plans to change rarely, such as code that estimates models, runs simulations, etc., should ideally be separated from fast code that one expects to change often, such as code that computes summary statistics, outputs tables, etc. This makes it easy to modify the presentation of the output without having to rerun the slow steps over and over.

Consider again a directory that estimates several specifications that take, together, an hour to run. The code that produces tables from those specifications will likely run in seconds. Therefore, it should be stored in a separate directory. We are likely to want to make many small changes to how we format the output, none of which will affect which specifications we want to run. It will not be efficient to have to repeatedly re-estimate the same model in order to change, say, the order of presentation of the parameters in the table.

Again, if instead estimation were very fast, it might be reasonable to include the script that produces tables inside the same directory as the script that estimates models. In such a case, the decision should be based on clarity, robustness, and the other principles articulated above, rather than on economizing CPU time.

For Further Reading

The ideas in this handbook are not new. The chapters are an attempt to communicate well-trod ideas from software engineering and computer science to a social science audience.

Here we list additional resources that have influenced our thinking. You may find these helpful if you wish to see some of the topics we have covered in greater depth.

Note that the single best resource we know of is not a book, but a website called software carpentry (http://software-carpentry.org/) devoted to teaching computing to scientists.

Bowman, Judith S., Sandra L. Emerson and Marcy Darnovsky. 2001. *The Practical SQL Handbook: Using SQL Variants.* New York: Addison-Wesley. [Chapter 2 contains a nice overview of database design.]

Brooks, Frederick P. 1975. *The Mythical Man-month.* Reading: Addison-Wesley.

Hunt, Andrew and David Thomas. 2000. *The Pragmatic Programmer: From Journeyman to Master.* Addison-Wesley: New York.

Immon, William H. 2005. *Building the Data Warehouse*. New York: Wiley.

Martin, Robert C. 2008. *Clean Code: A Handbook of Agile Software Craftsmanship.* New York: Prentice Hall.

Lutz, Mark and David Ascher. 1999. *Learning Python.* New York: O'Reilly Media.

Acknowledgments

We benefited greatly from the input of coauthors and colleagues on the methods described in this handbook. Ben Skrainka's Institute for Computational Economics lecture slides showed us that we were not alone.

Most importantly, we acknowledge the tireless work of the research assistants who suffered through our obsessions and wrong turns.